

crop-recommendation-system

July 9, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn import tree
from sklearn.metrics import accuracy_score
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_csv('/content/Crop_recommendation.csv')
df.head()
```

```
[2]:
```

	N	P	K	temperature	humidity	ph	rainfall	label
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice

```
[3]: df.tail()
```

```
[3]:
```

	N	P	K	temperature	humidity	ph	rainfall	label
2195	107	34	32	26.774637	66.413269	6.780064	177.774507	coffee
2196	99	15	27	27.417112	56.636362	6.086922	127.924610	coffee
2197	118	33	30	24.131797	67.225123	6.362608	173.322839	coffee
2198	117	32	34	26.272418	52.127394	6.758793	127.175293	coffee
2199	104	18	30	23.603016	60.396475	6.779833	140.937041	coffee

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2200 entries, 0 to 2199
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   N                2200 non-null   int64
```

```

1   P           2200 non-null   int64
2   K           2200 non-null   int64
3   temperature 2200 non-null   float64
4   humidity    2200 non-null   float64
5   ph          2200 non-null   float64
6   rainfall    2200 non-null   float64
7   label       2200 non-null   object
dtypes: float64(4), int64(3), object(1)
memory usage: 137.6+ KB

```

```
[5]: df.isnull().sum()
```

```

[5]: N           0
     P           0
     K           0
     temperature 0
     humidity    0
     ph          0
     rainfall    0
     label       0
     dtype: int64

```

```
[6]: df.size
```

```
[6]: 17600
```

```
[7]: df.shape
```

```
[7]: (2200, 8)
```

```
[8]: df.columns
```

```

[8]: Index(['N', 'P', 'K', 'temperature', 'humidity', 'ph', 'rainfall', 'label'],
          dtype='object')

```

```
[9]: df['label'].unique()
```

```

[9]: array(['rice', 'maize', 'chickpea', 'kidneybeans', 'pigeonpeas',
          'mothbeans', 'mungbean', 'blackgram', 'lentil', 'pomegranate',
          'banana', 'mango', 'grapes', 'watermelon', 'muskmelon', 'apple',
          'orange', 'papaya', 'coconut', 'cotton', 'jute', 'coffee'],
          dtype=object)

```

```
[10]: df.dtypes
```

```

[10]: N           int64
     P           int64

```

```

K                int64
temperature      float64
humidity         float64
ph              float64
rainfall        float64
label           object
dtype: object

```

```
[11]: df['label'].value_counts()
```

```

[11]: label
rice           100
maize          100
jute           100
cotton         100
coconut        100
papaya         100
orange         100
apple          100
muskmelon     100
watermelon    100
grapes         100
mango          100
banana         100
pomegranate   100
lentil         100
blackgram     100
mungbean      100
mothbeans     100
pigeonpeas    100
kidneybeans   100
chickpea      100
coffee        100
Name: count, dtype: int64

```

```
[33]: crop_summary = pd.pivot_table(df, index=['label'], aggfunc='mean')
crop_summary
```

```

[33]:
      K      N      P  humidity      ph  rainfall  \
label
apple   199.89  20.80  134.22  92.333383  5.929663  112.654779
banana   50.05  100.23   82.01  80.358123  5.983893  104.626980
blackgram  19.24   40.02   67.47  65.118426  7.133952   67.884151
chickpea  79.92   40.09   67.79  16.860439  7.336957   80.058977
coconut   30.59   21.98   16.93  94.844272  5.976562  175.686646
coffee   29.94  101.20   28.74  58.869846  6.790308  158.066295
cotton    19.56  117.77   46.24  79.843474  6.912675   80.398043

```


grapes	200.11	23.18	132.53	81.875228	6.025937	69.611829
jute	39.99	78.40	46.86	79.639864	6.732778	174.792798
kidneybeans	20.05	20.75	67.54	21.605357	5.749411	105.919778
lentil	19.41	18.77	68.36	64.804785	6.927932	45.680454
maize	19.79	77.76	48.44	65.092249	6.245190	84.766988
mango	29.92	20.07	27.18	50.156573	5.766373	94.704515
mothbeans	20.23	21.44	48.01	53.160418	6.831174	51.198487
mungbean	19.87	20.99	47.28	85.499975	6.723957	48.403601
muskmelon	50.08	100.32	17.72	92.342802	6.358805	24.689952
orange	10.01	19.58	16.55	92.170209	7.016957	110.474969
papaya	50.04	49.88	59.05	92.403388	6.741442	142.627839
pigeonpeas	20.29	20.73	67.73	48.061633	5.794175	149.457564
pomegranate	40.21	18.87	18.75	90.125504	6.429172	107.528442
rice	39.87	79.89	47.58	82.272822	6.425471	236.181114
watermelon	50.22	99.42	17.00	85.160375	6.495778	50.786219

	temperature
label	
apple	22.630942
banana	27.376798
blackgram	29.973340
chickpea	18.872847
coconut	27.409892
coffee	25.540477
cotton	23.988958
grapes	23.849575
jute	24.958376
kidneybeans	20.115085
lentil	24.509052
maize	22.389204
mango	31.208770
mothbeans	28.194920
mungbean	28.525775
muskmelon	28.663066
orange	22.765725
papaya	33.723859
pigeonpeas	27.741762
pomegranate	21.837842
rice	23.689332
watermelon	25.591767

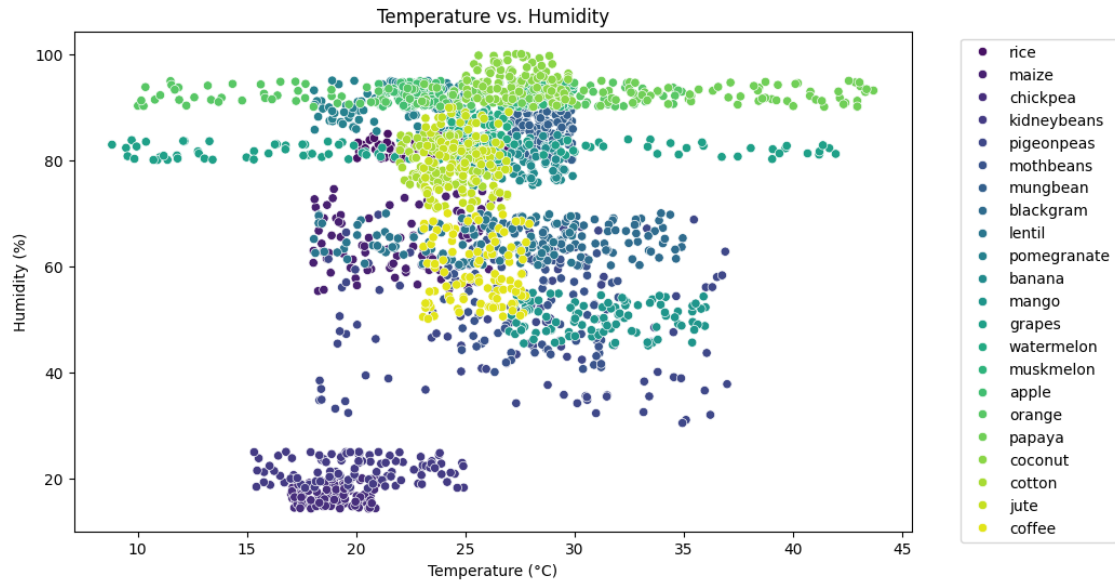
```
[12]: # Visualization 1: Histogram of Nitrogen (N) Content
plt.figure(figsize=(10, 6))
sns.histplot(df['N'], kde=True, color='blue')

plt.title('Distribution of Nitrogen Content in Soil')
plt.xlabel('Nitrogen Content (N)')
```

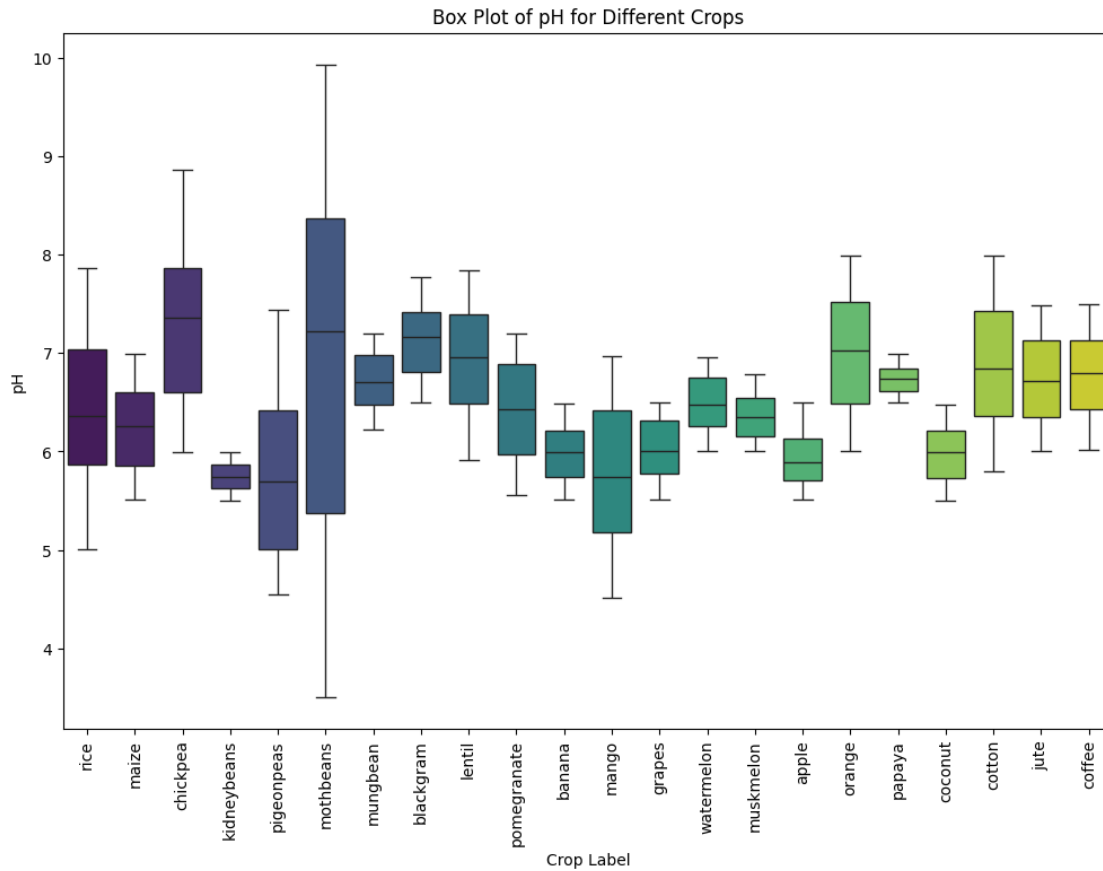
```
plt.ylabel('Frequency')
plt.show()
```



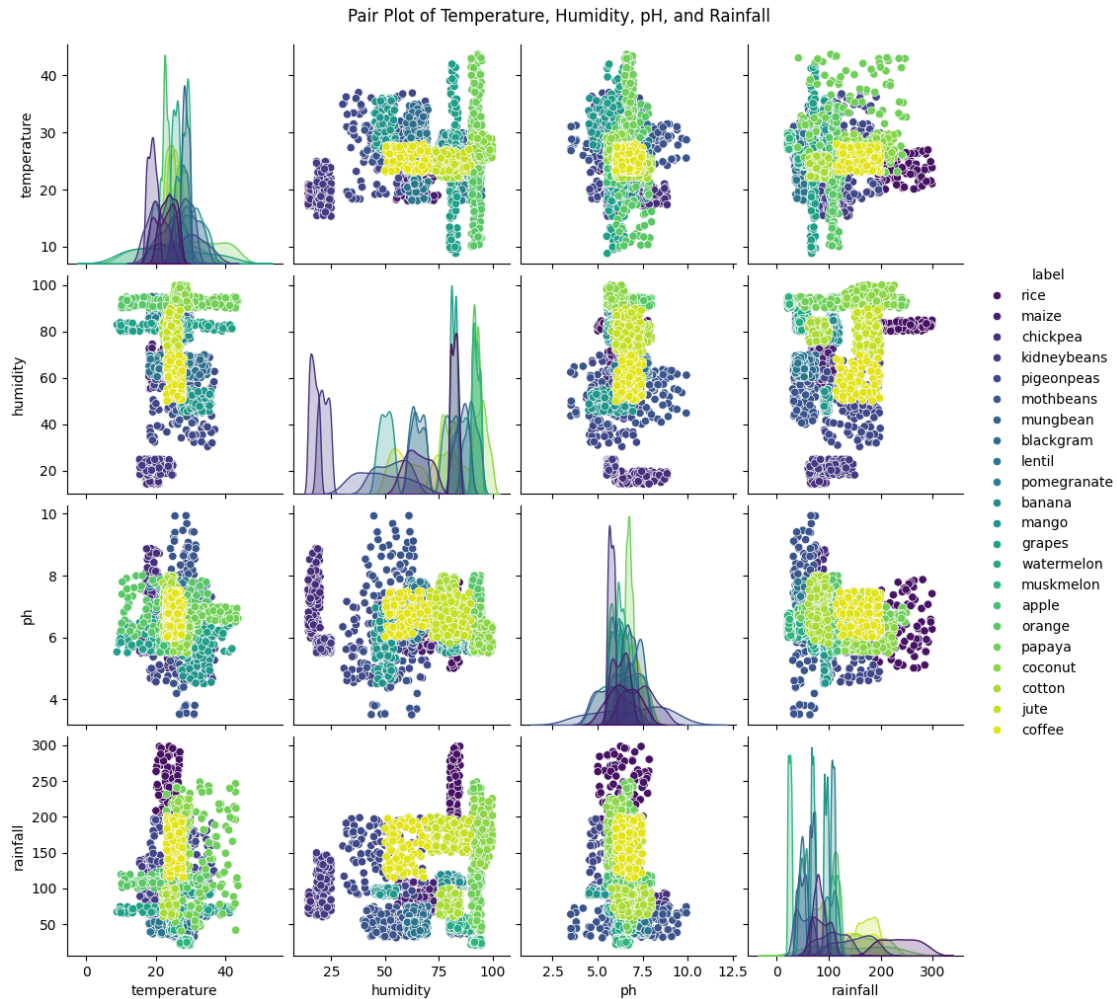
```
[13]: # Visualization 2: Scatter Plot of Temperature vs. Humidity
plt.figure(figsize=(10, 6))
sns.scatterplot(x=df['temperature'], y=df['humidity'], hue=df['label'],
               palette='viridis')
plt.title('Temperature vs. Humidity')
plt.xlabel('Temperature (°C)')
plt.ylabel('Humidity (%)')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()
```



```
[14]: # Visualization 3: Box Plot of pH for Different Crops
plt.figure(figsize=(12, 8))
sns.boxplot(x='label', y='ph', data=df, palette='viridis')
plt.title('Box Plot of pH for Different Crops')
plt.xlabel('Crop Label')
plt.ylabel('pH')
plt.xticks(rotation=90)
plt.show()
```




```
[15]: # Visualization 4: Pair Plot of Temperature, Humidity, pH, and Rainfall
pairplot_features = ['temperature', 'humidity', 'ph', 'rainfall', 'label']
sns.pairplot(df[pairplot_features], hue='label', palette='viridis',
             ↳diag_kind='kde')
plt.suptitle('Pair Plot of Temperature, Humidity, pH, and Rainfall', y=1.02)
plt.show()
```

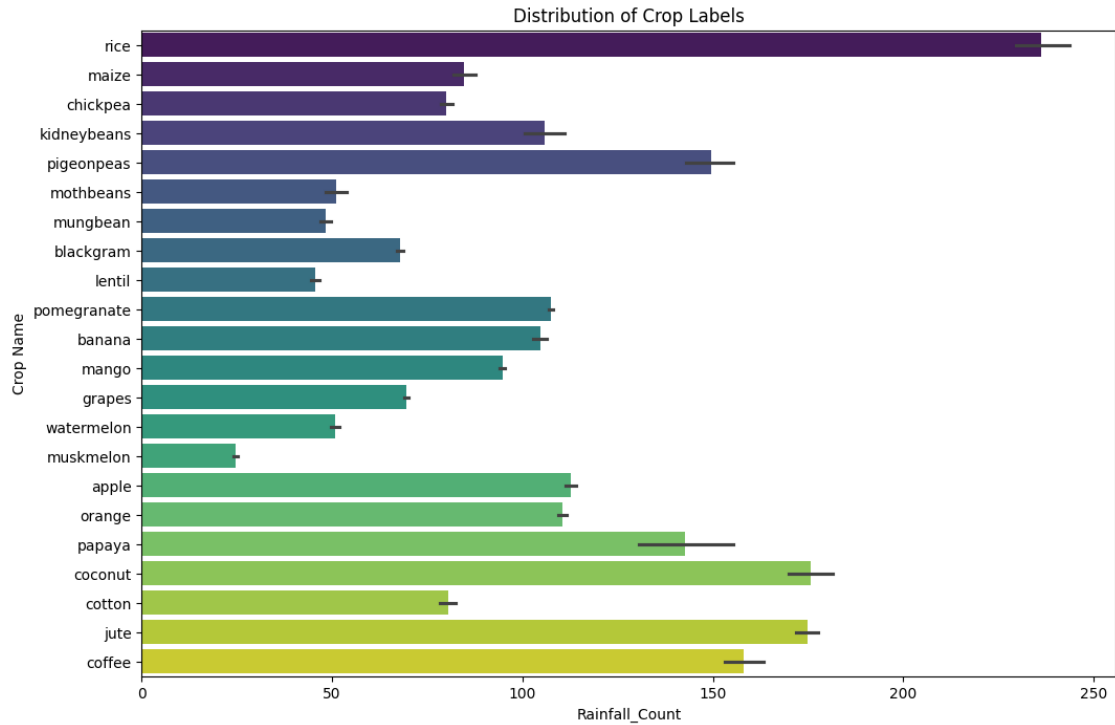


```
[16]: # Ensure only numerical columns are included in the correlation matrix
numerical_data = df.select_dtypes(include=['float64', 'int64'])
```

```
[17]: # Visualization 5: Heatmap of the Correlation Matrix
plt.figure(figsize=(10, 8))
sns.heatmap(numerical_data.corr(),annot=True,cmap='viridis', fmt='.2f')
plt.title('Heatmap of Correlation Matrix')
plt.show()
```

```
[18]: # Visualization: Bar Plot of Crop Label Counts
plt.figure(figsize=(12, 8))
sns.barplot(y='label', x='rainfall', data=df, palette='viridis')
plt.title('Distribution of Crop Labels')
plt.xlabel('Rainfall_Count')
plt.ylabel('Crop Name')
plt.show()
```



```
[19]: x = df.drop('label', axis = 1)
      y = df['label']
```

```
[21]: from sklearn.model_selection import train_test_split
      x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.
      ↪2,random_state =2)
```

#1.Decision Tree

```
[22]: from sklearn.tree import DecisionTreeClassifier
      model_1 = DecisionTreeClassifier(criterion='entropy',max_depth = 6,
      ↪random_state = 2)
      model_1.fit(x_train, y_train)
      y_pred_1 = model_1.predict(x_test)
      decision_acc = accuracy_score(y_test, y_pred_1)
      print("Accuracy of decision tree is : = " + str(decision_acc))
      print(classification_report(y_test,y_pred_1))
```

Accuracy of decision tree is : = 0.9590909090909091

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	0.80	1.00	0.89	16

chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	1.00	1.00	1.00	20
grapes	1.00	1.00	1.00	18
jute	0.84	0.96	0.90	28
kidneybeans	1.00	0.79	0.88	14
lentil	0.95	0.83	0.88	23
maize	1.00	1.00	1.00	21
mango	1.00	1.00	1.00	26
mothbeans	0.67	0.74	0.70	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	0.92	0.69	0.79	16
watermelon	1.00	1.00	1.00	15
accuracy			0.96	440
macro avg	0.96	0.95	0.96	440
weighted avg	0.96	0.96	0.96	440

##2.LogisticRegression

```
[23]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
logistic_acc = accuracy_score(y_test, y_pred)
print("Accuracy of logistic regression is : = " + str(logistic_acc))
print(classification_report(y_test,y_pred))
```

Accuracy of logistic regression is : = 0.9522727272727273

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	0.86	0.75	0.80	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	0.86	0.90	0.88	20
grapes	1.00	1.00	1.00	18
jute	0.84	0.93	0.88	28
kidneybeans	1.00	1.00	1.00	14

lentil	0.88	1.00	0.94	23
maize	0.90	0.86	0.88	21
mango	0.96	1.00	0.98	26
mothbeans	0.84	0.84	0.84	19
mungbean	1.00	0.96	0.98	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	0.95	0.97	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	0.85	0.69	0.76	16
watermelon	1.00	1.00	1.00	15
accuracy			0.95	440
macro avg	0.95	0.95	0.95	440
weighted avg	0.95	0.95	0.95	440

##3.GaussianNB

```
[24]: from sklearn.naive_bayes import GaussianNB
model_3 = GaussianNB()
model_3.fit(x_train, y_train)
y_pred_3 = model_3.predict(x_test)
naive_bayes_acc = accuracy_score(y_test, y_pred_3)
print("Accuracy of naive_bayes is : = " + str(naive_bayes_acc))
print(classification_report(y_test,y_pred_3))
```

Accuracy of naive_bayes is : = 0.990909090909091

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	1.00	1.00	1.00	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	1.00	1.00	1.00	20
grapes	1.00	1.00	1.00	18
jute	0.88	1.00	0.93	28
kidneybeans	1.00	1.00	1.00	14
lentil	1.00	1.00	1.00	23
maize	1.00	1.00	1.00	21
mango	1.00	1.00	1.00	26
mothbeans	1.00	1.00	1.00	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29

papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	1.00	0.75	0.86	16
watermelon	1.00	1.00	1.00	15
accuracy			0.99	440
macro avg	0.99	0.99	0.99	440
weighted avg	0.99	0.99	0.99	440

##4.Support Vector Machine

```
[25]: from sklearn.svm import SVC
model_svm = SVC()
model_svm.fit(x_train, y_train)
y_pred_svm = model_svm.predict(x_test)
svm_acc = accuracy_score(y_test, y_pred_svm)
print("Accuracy of Support Vector Machine is: " + str(svm_acc))
print(classification_report(y_test,y_pred_svm))
```

Accuracy of Support Vector Machine is: 0.9772727272727273

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	0.94	1.00	0.97	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	0.95	1.00	0.98	20
grapes	1.00	1.00	1.00	18
jute	0.85	1.00	0.92	28
kidneybeans	0.93	1.00	0.97	14
lentil	0.92	1.00	0.96	23
maize	1.00	0.95	0.98	21
mango	1.00	1.00	1.00	26
mothbeans	1.00	0.84	0.91	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	0.94	0.97	18
pomegranate	1.00	1.00	1.00	17
rice	1.00	0.69	0.81	16
watermelon	1.00	1.00	1.00	15
accuracy			0.98	440

macro avg	0.98	0.97	0.98	440
weighted avg	0.98	0.98	0.98	440

##5.Random Forest

```
[26]: from sklearn.ensemble import RandomForestClassifier
model_rf = RandomForestClassifier()
model_rf.fit(x_train, y_train)
y_pred_rf = model_rf.predict(x_test)
rf_acc = accuracy_score(y_test, y_pred_rf)
print("Accuracy of Support Vector Machine is: " + str(rf_acc))
print(classification_report(y_test, y_pred_rf))
```

Accuracy of Support Vector Machine is: 0.9772727272727273

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	1.00	1.00	1.00	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	1.00	1.00	1.00	20
grapes	1.00	1.00	1.00	18
jute	0.93	1.00	0.97	28
kidneybeans	1.00	1.00	1.00	14
lentil	1.00	1.00	1.00	23
maize	1.00	1.00	1.00	21
mango	1.00	1.00	1.00	26
mothbeans	1.00	1.00	1.00	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	1.00	0.88	0.93	16
watermelon	1.00	1.00	1.00	15
accuracy			1.00	440
macro avg	1.00	0.99	1.00	440
weighted avg	1.00	1.00	1.00	440

[]:

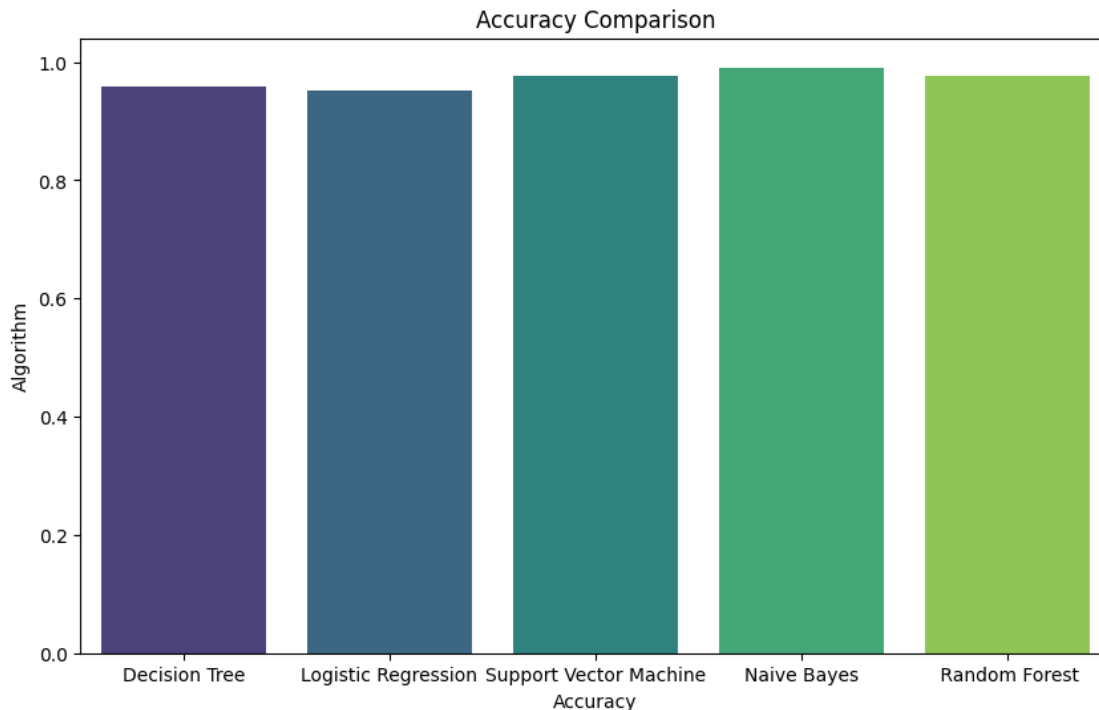
```
[27]: print("Accuracy of decision tree is : = " + str(decision_acc))
      print("Accuracy of logistic regression is : = " + str(logistic_acc))
      print("Accuracy of Support Vector Machine is: " + str(svm_acc))
      print("Accuracy of naive_bayes is : = " + str(naive_bayes_acc))
```

```
Accuracy of decision tree is : = 0.9590909090909091
Accuracy of logistic regression is : = 0.9522727272727273
Accuracy of Support Vector Machine is: 0.9772727272727273
Accuracy of naive_bayes is : = 0.9909090909090909
```

Accuracy Comparison

```
[28]: accuracies = [decision_acc, logistic_acc, svm_acc, naive_bayes_acc, rf_acc]
      models = ['Decision Tree', 'Logistic Regression', 'Support Vector Machine', 'Naive Bayes', 'Random Forest']
```

```
[29]: plt.figure(figsize=(10,6))
      sns.barplot(x=models, y=accuracies, palette='viridis')
      plt.title('Accuracy Comparison')
      plt.xlabel('Accuracy')
      plt.ylabel('Algorithm')
      plt.show()
```



0.1 Making a prediction

```
[30]: data = np.array([[104,18, 30, 23.603016, 60.3, 6.7, 140.91]])  
      prediction = model_3.predict(data)  
      print(prediction)
```

```
['coffee']
```

```
[31]: data = np.array([[60,      55,      44      ,23.004459,82.320763,7.  
      ↪840207,      263.964248]])  
      prediction = model_rf.predict(data)  
      print(prediction)
```

```
['rice']
```

Conclusion

The SVM model achieved the highest accuracy for the crop recommendation system, of approximately 0.18. This suggests that the SVM model is highly effective in predicting the appropriate crops based on soil and environmental features. For further improvements, ensemble techniques like Bagging and Boosting could be employed to enhance the model's performance. This system can significantly assist farmers in optimizing crop selection, leading to improved agricultural productivity and sustainability.

```
[ ]:
```